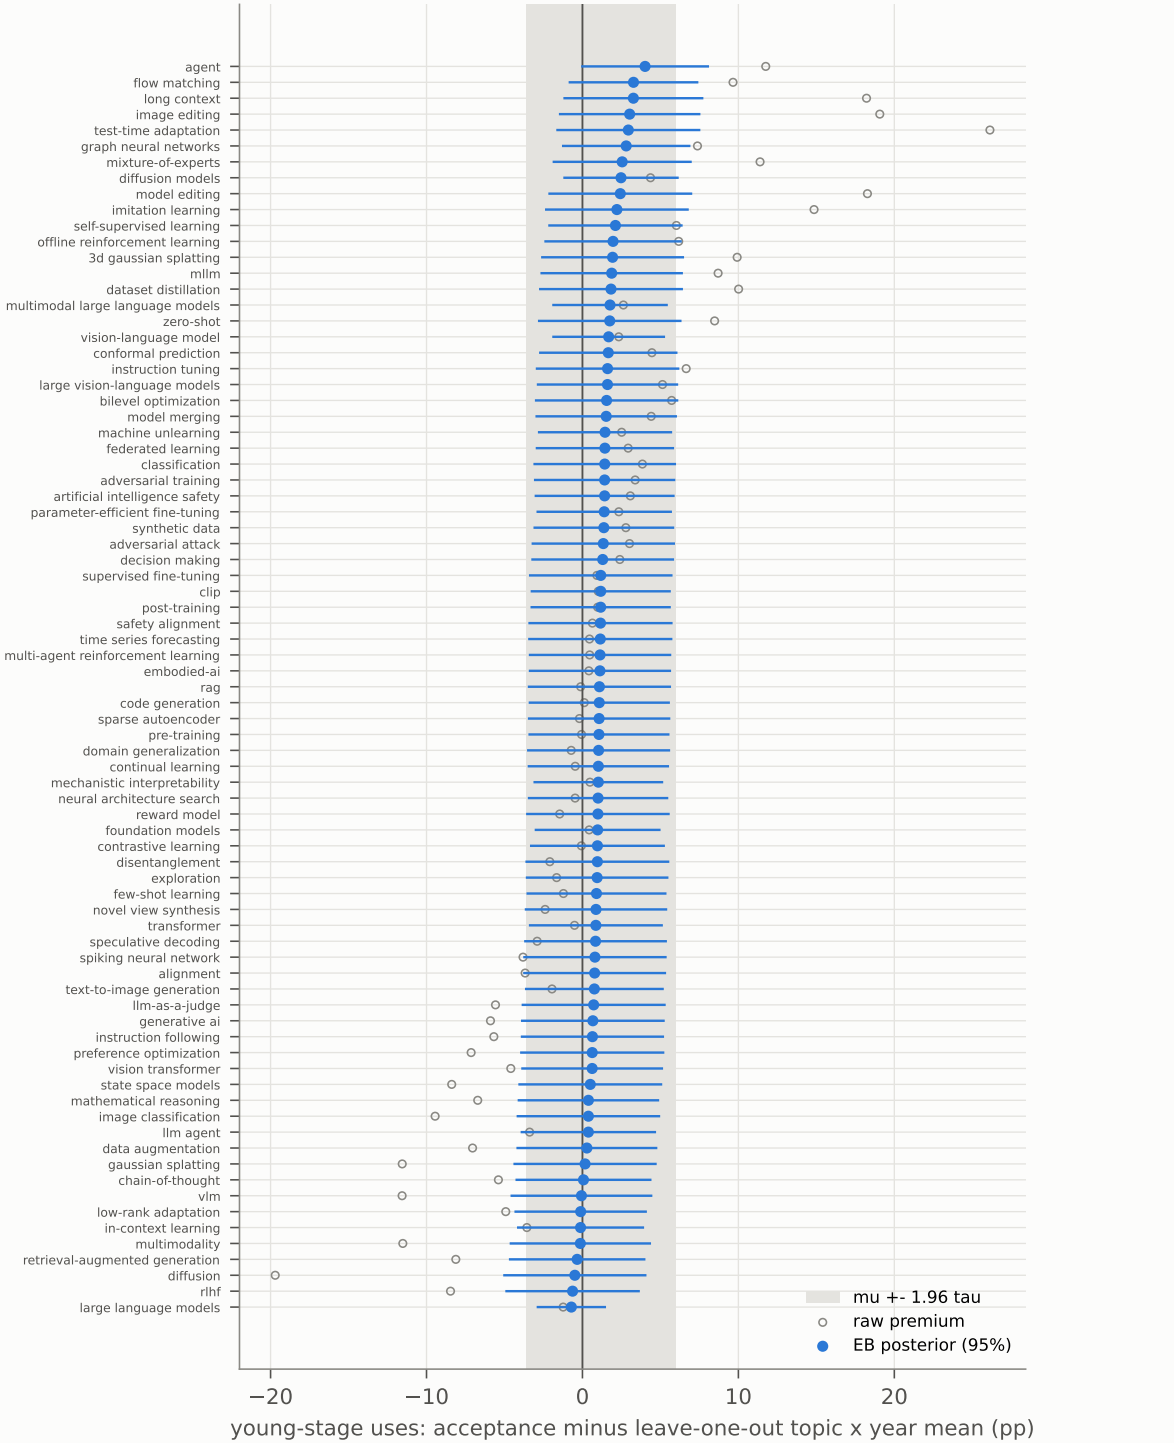
Early-entry premium by keyword: tau = 2.4 pp (permutation null mean 0.2, p = 0.04); 79 keywords
agent
flow matching
long context
image editing
test-time adaptation
graph neural networks
mixture-of-experts
diffusion models
model editing
imitation learning
self-supervised learning
offline reinforcement learning
3d gaussian splatting
mllm
dataset distillation
multimodal large language models
zero-shot
vision-language model
conformal prediction
instruction tuning
large vision-language models
bilevel optimization
model merging
machine unlearning
federated learning
classification
adversarial training
artificial intelligence safety
parameter-efficient fine-tuning
synthetic data
adversarial attack
decision making
supervised fine-tuning
clip
post-training
safety alignment
time series forecasting
multi-agent reinforcement learning
embodied-ai
rag
code generation
sparse autoencoder
pre-training
domain generalization
continual learning
mechanistic interpretability
neural architecture search
reward model
foundation models
contrastive learning
disentanglement
exploration
few-shot learning
novel view synthesis
transformer
speculative decoding
spiking neural network
alignment
text-to-image generation
llm-as-a-judge
generative ai
instruction following
preference optimization
vision transformer
state space models
mathematical reasoning
image classification
llm agent
data augmentation
gaussian splatting
chain-of-thought
vlm
low-rank adaptation
in-context learning
multimodality
retrieval-augmented generation
diffusion
rlhf
large language models
mu +- 1.96 tau
raw premium
EB posterior (95%)
−20
−10
0
10
20
young-stage uses: acceptance minus leave-one-out topic x year mean (pp)